

# Neural networks and Backpropagation

```
$ echo "Data Sciences Institute"
```

# Yesterday

- Overview
- Computation graph view of neural networks
- Linear operation followed by non-linear activation
  - ...But what is the linear operation, really?

# Today

- A closer look at what's going on in a "neuron"
- Backpropagation: how do we train a neural network?

# Neural Network for classification

Vector function with tunable parameters  $\theta$

$$\mathbf{f}(\cdot; \theta) : \mathbb{R}^N \rightarrow (0, 1)^K$$

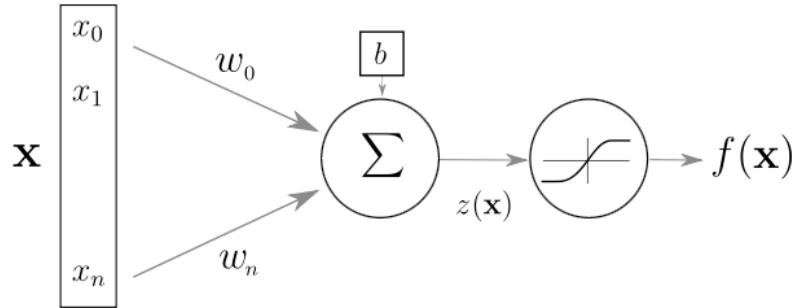
Sample  $s$  in dataset  $S$ :

- input:  $\mathbf{x}^s \in \mathbb{R}^N$
- expected output:  $y^s \in [0, K - 1]$

Output is a conditional probability distribution:

$$\mathbf{f}(\mathbf{x}^s; \theta)_c = P(Y = c | X = \mathbf{x}^s)$$

# Artificial Neuron



$$z(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

$$f(\mathbf{x}) = g(\mathbf{w}^T \mathbf{x} + b)$$

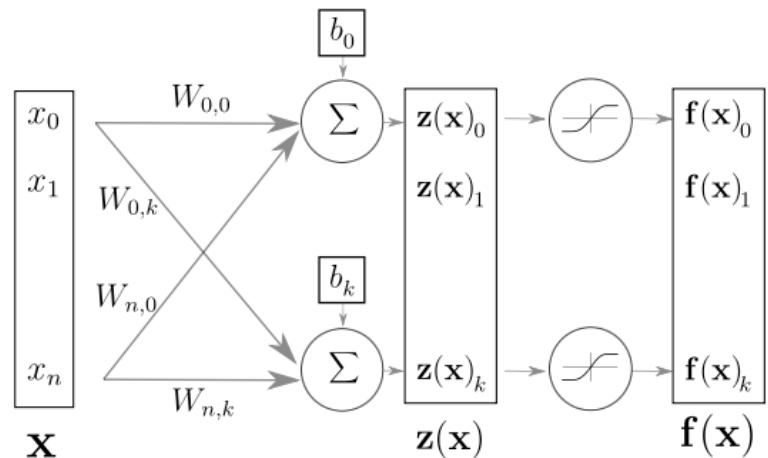
- $\mathbf{x}, f(\mathbf{x})$  input and output
- $z(\mathbf{x})$  pre-activation
- $\mathbf{w}, b$  weights and bias
- $g$  activation function



# Concrete Example

- Say we have two input dimensions  $x_1$  and  $x_2$  and one output dimension  $f(x)$  (sometimes,  $\hat{y}$  - the predicted value of  $y$  - is used instead of  $f(x)$ )
- Our weights and biases could be  $W = [3, -2]$  and  $b = 1$
- Our non-linearity could be ReLU:  $g(z) = \max(0, z)$
- Now  $z(x) = 3x_1 - 2x_2 + 1$  and  $f(x) = \max(0, 3x_1 - 2x_2 + 1)$
- Every neuron in a neural network is a function like this!

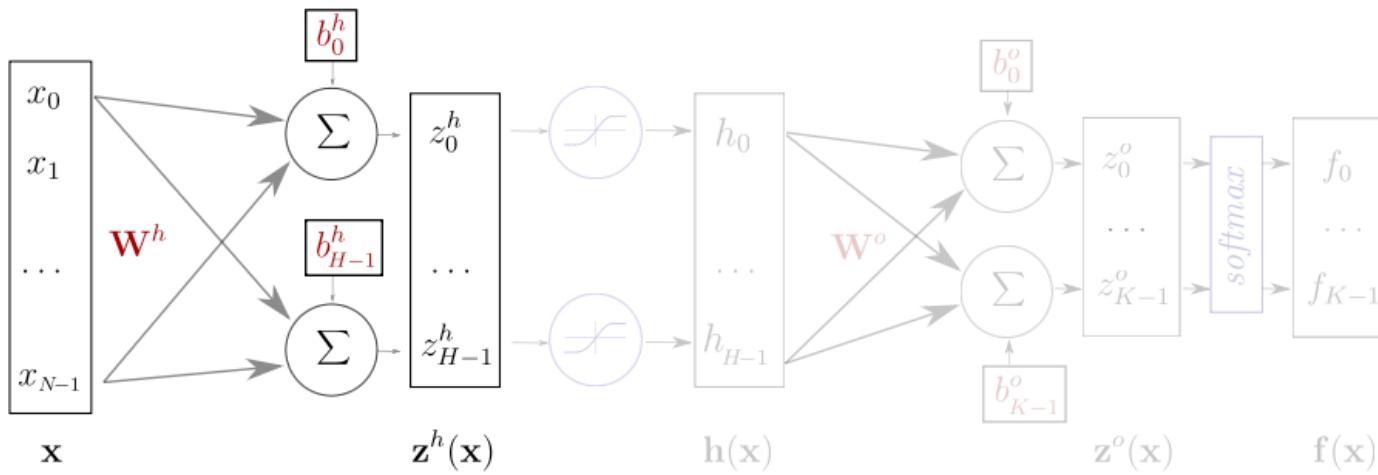
# Layer of Neurons (Vectorization)



$$\mathbf{f}(\mathbf{x}) = g(\mathbf{z}(\mathbf{x})) = g(\mathbf{W}\mathbf{x} + \mathbf{b})$$

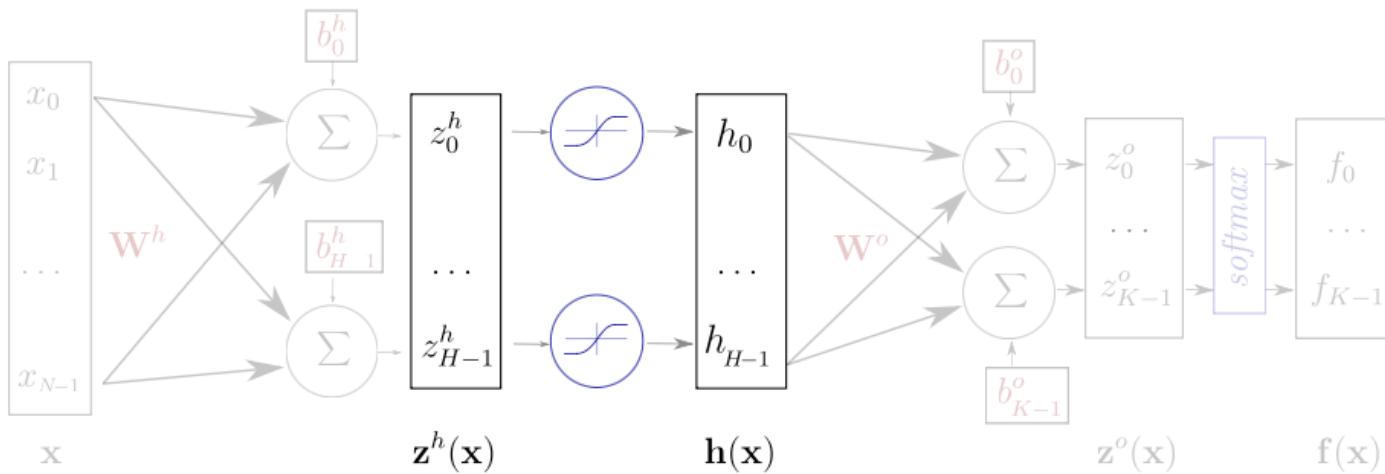
- $\mathbf{W}, \mathbf{b}$  now matrix and vector

# One Hidden Layer Network



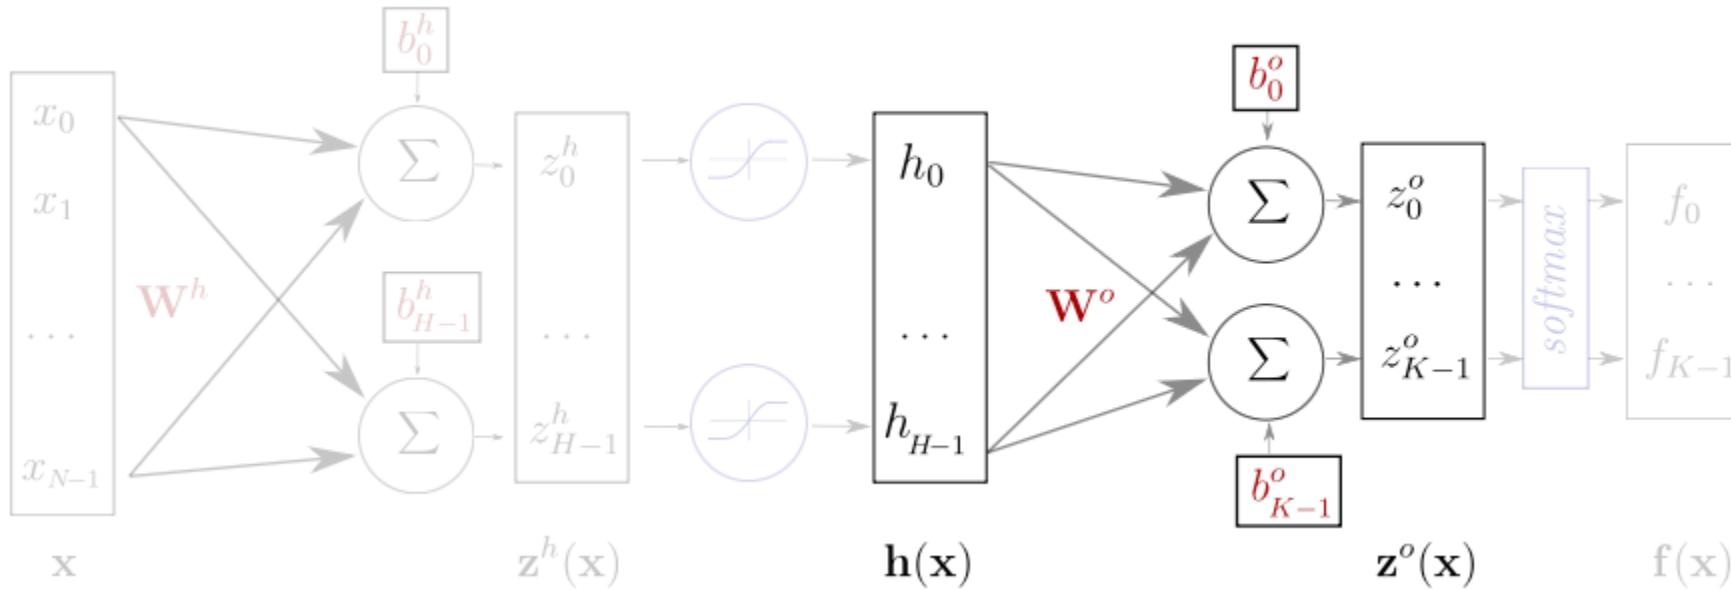
- $\mathbf{z}^h(\mathbf{x}) = \mathbf{W}^h \mathbf{x} + \mathbf{b}^h$

# One Hidden Layer Network



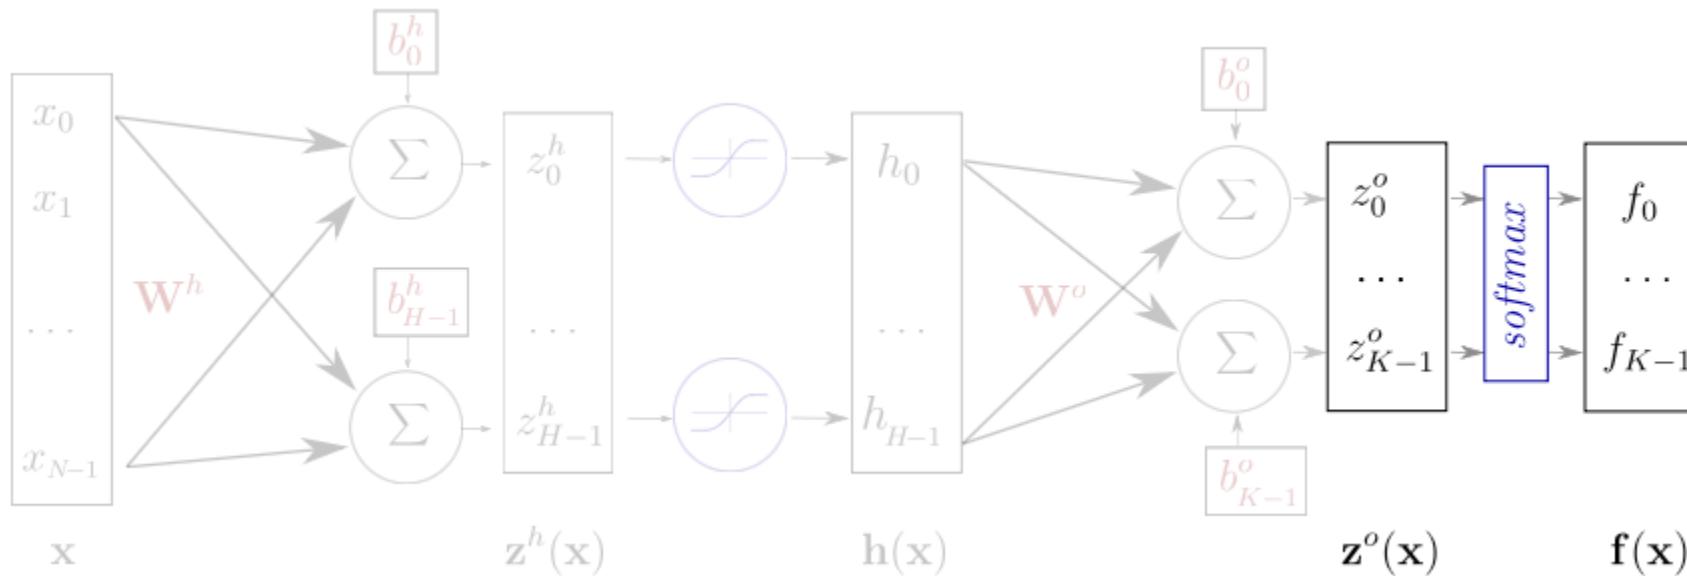
- $\mathbf{h}(\mathbf{x}) = g(\mathbf{z}^h(\mathbf{x})) = g(\mathbf{W}^h \mathbf{x})$

# One Hidden Layer Network



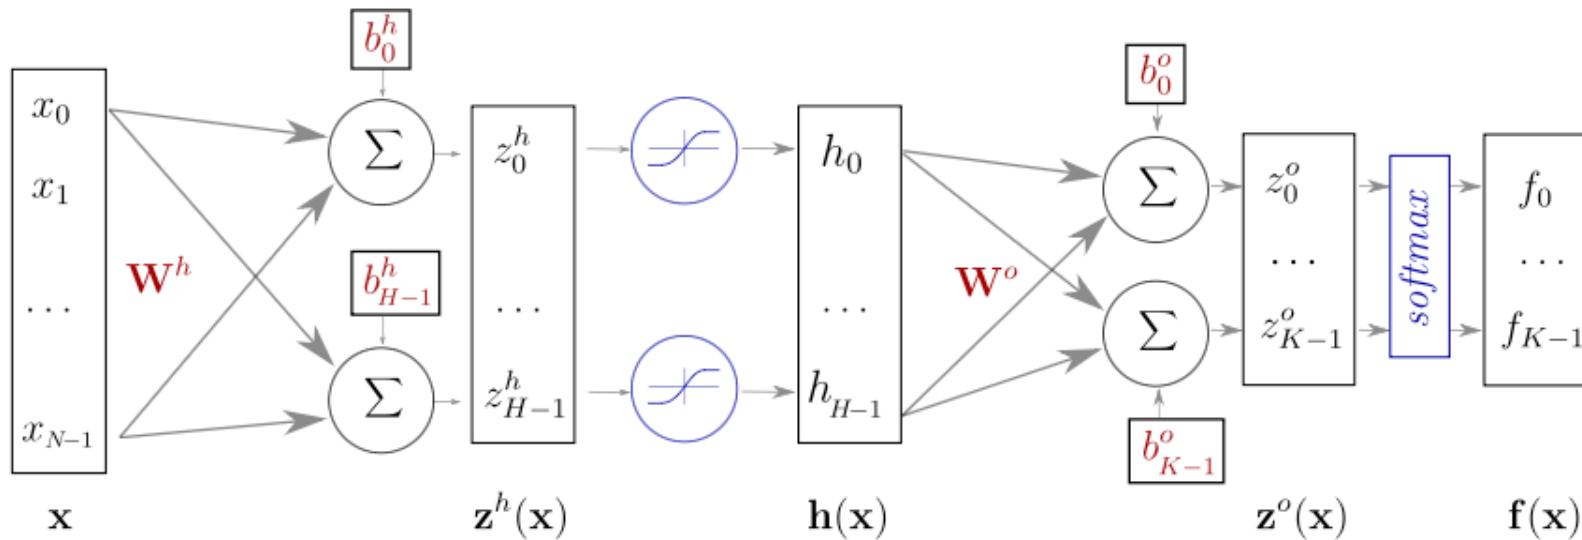
- $\mathbf{z}^o(\mathbf{x}) = \mathbf{W}^o\mathbf{h}(\mathbf{x}) + \mathbf{b}^o$

# One Hidden Layer Network

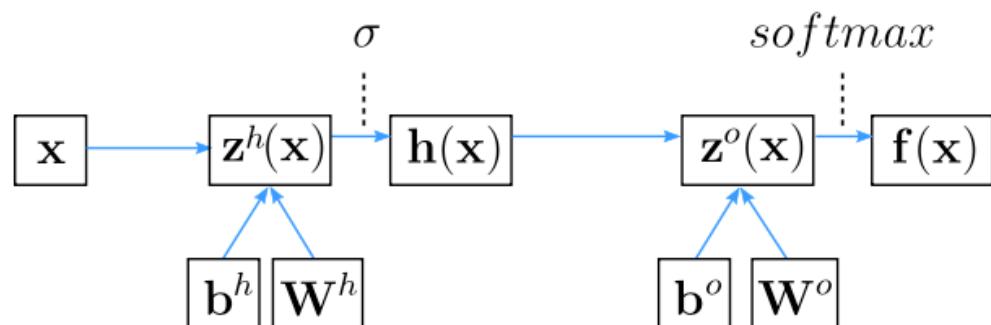


- $\mathbf{f}(\mathbf{x}) = softmax(\mathbf{z}^o) = softmax(\mathbf{W}^o\mathbf{h}(\mathbf{x}) + \mathbf{b}^o)$

# One Hidden Layer Network



## Alternate representation

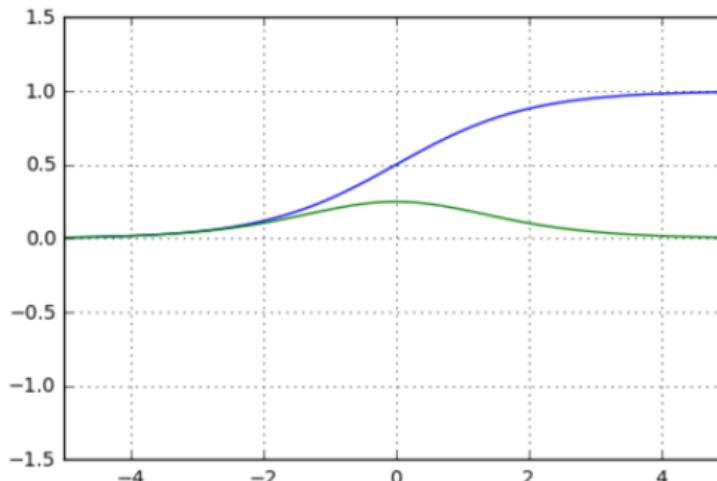


# One Hidden Layer Network

## Keras implementation

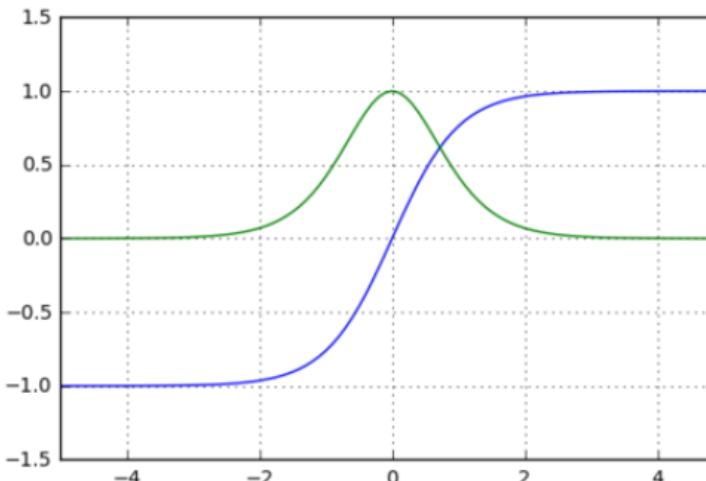
```
model = Sequential()
model.add(Dense(H, input_dim=N))      # weight matrix dim [N * H]
model.add(Activation("tanh"))
model.add(Dense(K))                  # weight matrix dim [H x K]
model.add(Activation("softmax"))
```

# Element-wise activation functions



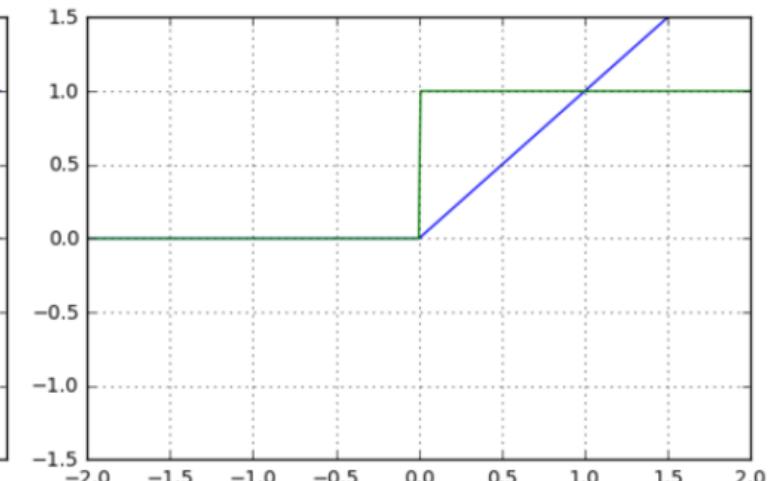
$$\text{sigm}(x) = \frac{1}{1 + e^{-x}}$$

$$\text{sigm}'(x) = \text{sigm}(x)(1 - \text{sigm}(x))$$



$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$

$$\tanh'(x) = 1 - \tanh(x)^2$$



$$\text{relu}(x) = \max(0, x)$$

$$\text{relu}'(x) = 1_{x>0}$$

- blue: activation function
- green: derivative



# Softmax function

$$\text{softmax}(\mathbf{x}) = \frac{1}{\sum_{i=1}^n e^{x_i}} \cdot \begin{bmatrix} e^{x_1} \\ e^{x_2} \\ \vdots \\ e^{x_n} \end{bmatrix}$$

- vector of values in  $(0, 1)$  that add up to 1
- for example,  $\mathbf{x} = [1, 2, 3]$  becomes  $\frac{1}{e^1 + e^2 + e^3} \cdot [e^1, e^2, e^3]^T = [0.09, 0.24, 0.67]$
- $p(Y = c | X = \mathbf{x}) = \text{softmax}(\mathbf{z}(\mathbf{x}))_c$
- the pre-activation vector  $\mathbf{z}(\mathbf{x})$  is often called "the logits"

# Training the network

Find parameters that minimize the **negative log likelihood** (or cross entropy)

The loss function for a given sample  $s \in S$ :

$$l(\mathbf{f}(\mathbf{x}^s; \theta), y^s) = nll(\mathbf{x}^s, y^s; \theta) = -\log \mathbf{f}(\mathbf{x}^s; \theta)_{y^s}$$

**example**  $y^s = 3$

$$l(\mathbf{f}(\mathbf{x}^s; \theta), y^s) = l \left( \begin{array}{c|c} \begin{matrix} f_0 \\ \dots \\ f_3 \\ \dots \\ f_{K-1} \end{matrix}, & \begin{matrix} 0 \\ \dots \\ 1 \\ \dots \\ 0 \end{matrix} \end{array} \right) = -\log f_3$$



# Training the network

Find parameters  $\theta = (\mathbf{W}^h; \mathbf{b}^h; \mathbf{W}^o; \mathbf{b}^o)$  that minimize the **negative log likelihood** (or [cross entropy](#))

The loss function for a given sample  $s \in S$ :

$$l(\mathbf{f}(\mathbf{x}^s; \theta), y^s) = nll(\mathbf{x}^s, y^s; \theta) = -\log \mathbf{f}(\mathbf{x}^s; \theta)_{y^s}$$

The cost function is the negative likelihood of the model computed on the full training set (for i.i.d. samples):

$$L_S(\theta) = -\frac{1}{|S|} \sum_{s \in S} \log \mathbf{f}(\mathbf{x}^s; \theta)_{y^s}$$



# Training the network

- Now we have a mathematical function representing the network
- And we have a way of measuring how good it is
- How do we find the parameters that minimize the loss?

# Gradient Descent

- Let's imagine we only have one parameter  $\theta$
- We can compute the derivative of the loss with respect to  $\theta$
- The derivative,  $\frac{dL}{d\theta}$ , tells us the slope of the loss function at a given point
- If  $\frac{dL}{d\theta} > 0$ , increasing  $\theta$  will increase the loss, and vice versa
- To minimize loss, we adjust  $\theta$  in the opposite direction of  $\frac{dL}{d\theta}$
- This is done using the update rule:  $\theta = \theta_{old} - \eta \frac{dL}{d\theta}$

# Gradient Descent

- We can use gradient descent to play "guess what number I'm thinking of"
- If your guess is too high, you decrease it
- If your guess is too low, you increase it
- The error function is a parabola
- By finding the lowest point on the parabola, you find the best guess

# Implementing Gradient Descent

- Start with an initial guess for  $\theta$
- Calculate  $\frac{dL}{d\theta}$  using the current value of  $\theta$
- Update  $\theta$  using the update rule
- Repeat the process until the change in loss is below a threshold or a set number of iterations is reached
- The choice of learning rate  $\eta$  is crucial: too high, and we may overshoot the minimum; too low, and convergence will be slow

# Stochastic Gradient Descent

- Traditional Gradient Descent uses the entire dataset to compute the gradient, which can be computationally expensive
- Stochastic Gradient Descent (SGD) updates the parameters using only a single data point (or a small batch)
- In SGD, for each iteration, a data point (or batch) is randomly selected to compute the gradient
- Since only a subset of data is used, the gradient estimation can be noisy, leading to a less smooth path towards the minimum
- However, SGD is much faster than traditional gradient descent

# Stochastic Gradient Descent

Initialize  $\theta$  randomly

For  $E$  epochs perform:

- Randomly select a small batch of samples ( $B \subset S$ )
  - Compute gradients:  $\Delta = \nabla_{\theta} L_B(\theta)$
  - Update parameters:  $\theta \leftarrow \theta - \eta \Delta$
- Repeat until the epoch is completed (all of  $S$  is covered)  
Stop when reaching criterion:
- $\text{nll}$  stops decreasing when computed on validation set



# Computing Gradients

Output Weights:  $\frac{\partial l(\mathbf{f}(\mathbf{x}), y)}{\partial W_{i,j}^o}$

Hidden Weights:  $\frac{\partial l(\mathbf{f}(\mathbf{x}), y)}{\partial W_{i,j}^h}$

Output bias:  $\frac{\partial l(\mathbf{f}(\mathbf{x}), y)}{\partial b_i^o}$

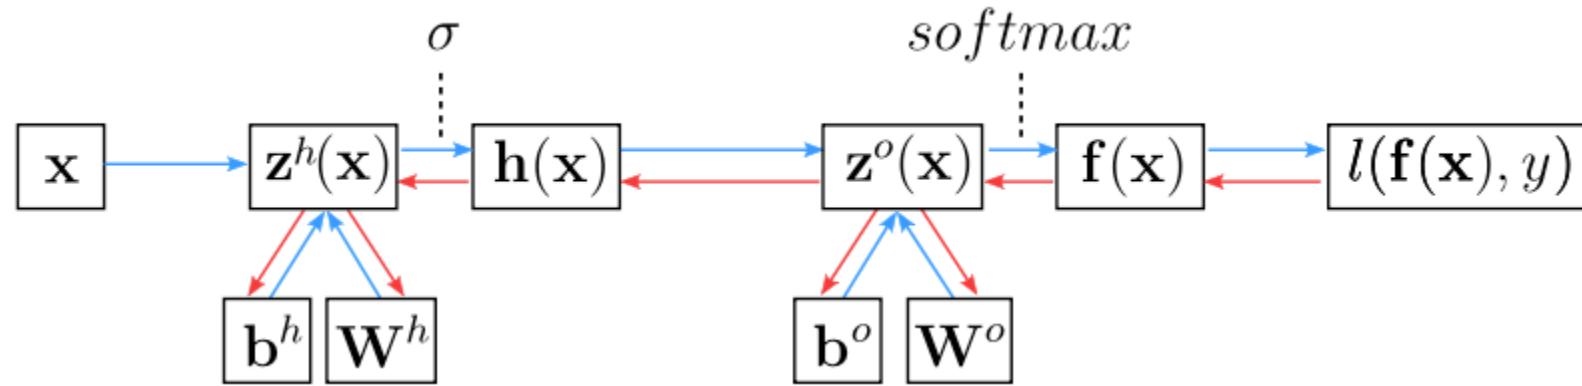
Hidden bias:  $\frac{\partial l(\mathbf{f}(\mathbf{x}), y)}{\partial b_i^h}$

- The network is a composition of differentiable modules
- We can apply the "chain rule"

# Chain rule

- Mathematical theorem that lets us compute derivatives when functions are inside other functions
- Remember, our neural network is a composition of functions:  $f(x) = g(h(x))$
- The chain rule tells us how to compute  $\frac{df}{dx}$
- $\frac{df}{dx} = \frac{df}{dg} \frac{dg}{dh} \frac{dh}{dx}$
- In English: The derivative of the overall network with respect to its input is the product of derivatives of each function in the network

# Backpropagation



- Compute partial derivatives of the loss
- For any given function in the network, we can compute how changing its parameters will affect the loss
- In other words, we can find how much each parameter's value contributes to the loss

# Initialization and Learning Tricks

# Initialization and normalization

- Input data should be normalized to have approx. same range:
  - standardization or quantile normalization
- Initializing weights:
  - Zero is a saddle point: no gradient, no learning
  - Constant init: all neurons compute the same function
  - Solution: random init, ex:  $w \sim \mathcal{N}(0, 0.1)$
  - Better inits: Xavier Glorot and Kaming He & orthogonal
- Biases can (should) be initialized to zero

# SGD learning rate

- Very sensitive:
  - Too high → early plateau or even divergence
  - Too low → slow convergence
  - Try a large value first:  $\eta = 0.1$  or even  $\eta = 1$
  - Divide by 10 and retry in case of divergence
- Large constant LR prevents final convergence
  - multiply  $\eta_t$  by  $\beta < 1$  after each update
  - or monitor validation loss and divide  $\eta_t$  by 2 or 10 when no progress
  - See [ReduceLROnPlateau](#) in Keras

# Momentum

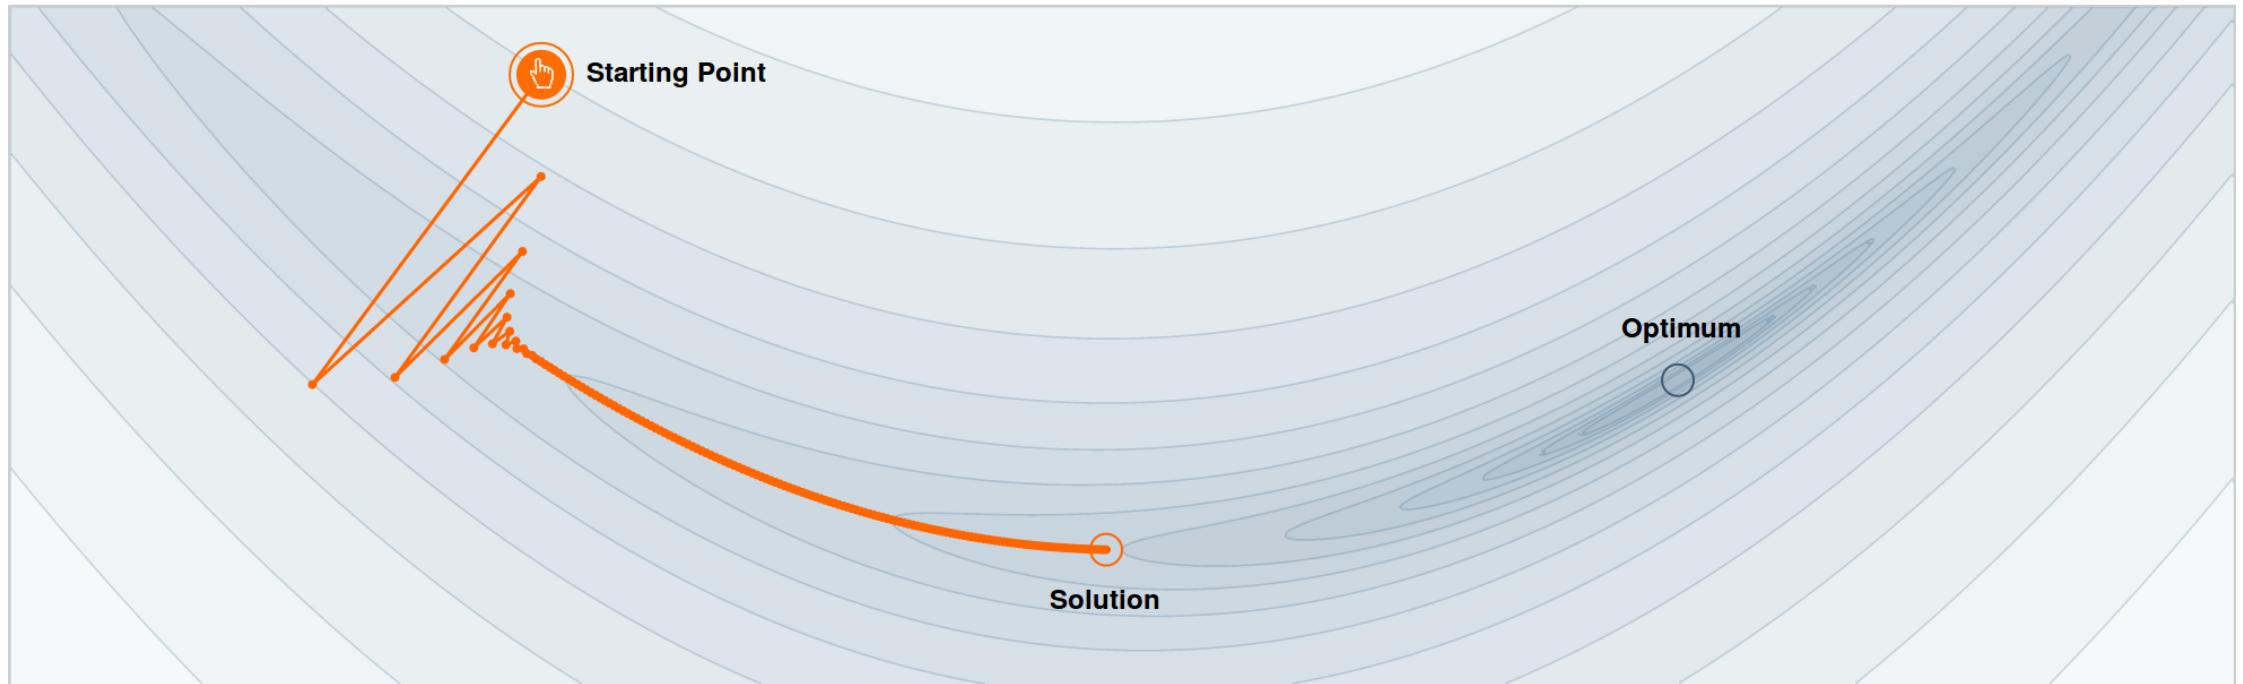
Accumulate gradients across successive updates:

$$m_t = \gamma m_{t-1} + \eta \nabla_{\theta} L_{B_t}(\theta_{t-1})$$

$$\theta_t = \theta_{t-1} - m_t$$

$\gamma$  is typically set to 0.9

Larger updates in directions where the gradient sign is constant  
to accelerate in low curvature areas



Step-size  $\alpha = 0.0030$

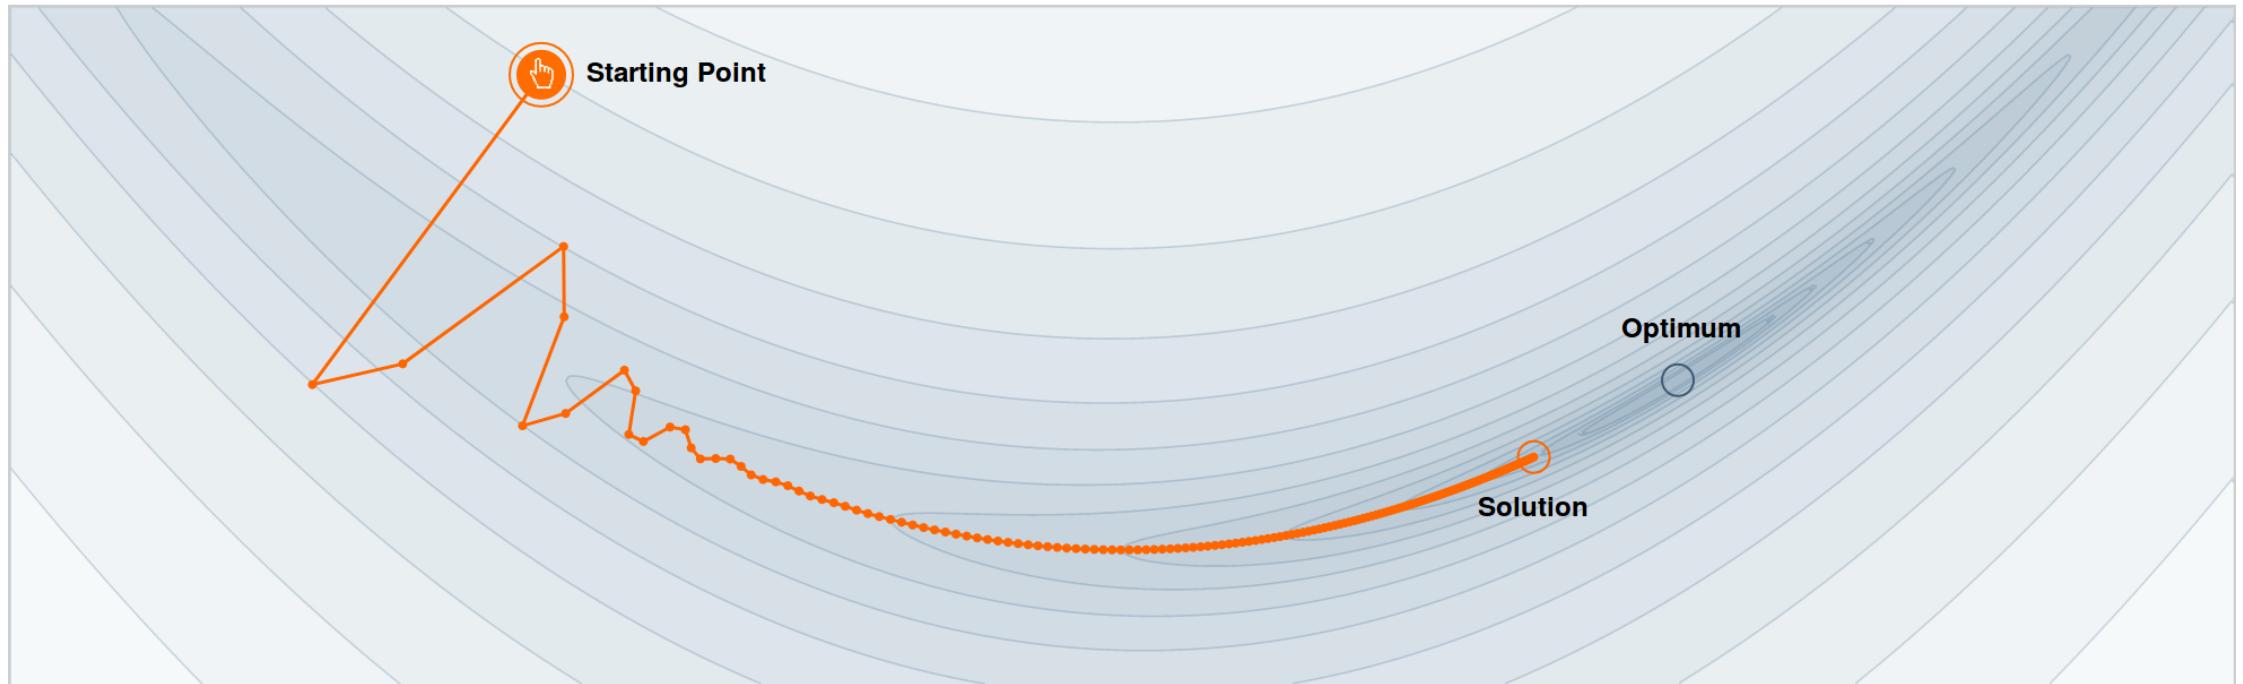


Momentum  $\beta = 0.0$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

## Why Momentum Really Works



Step-size  $\alpha = 0.0030$

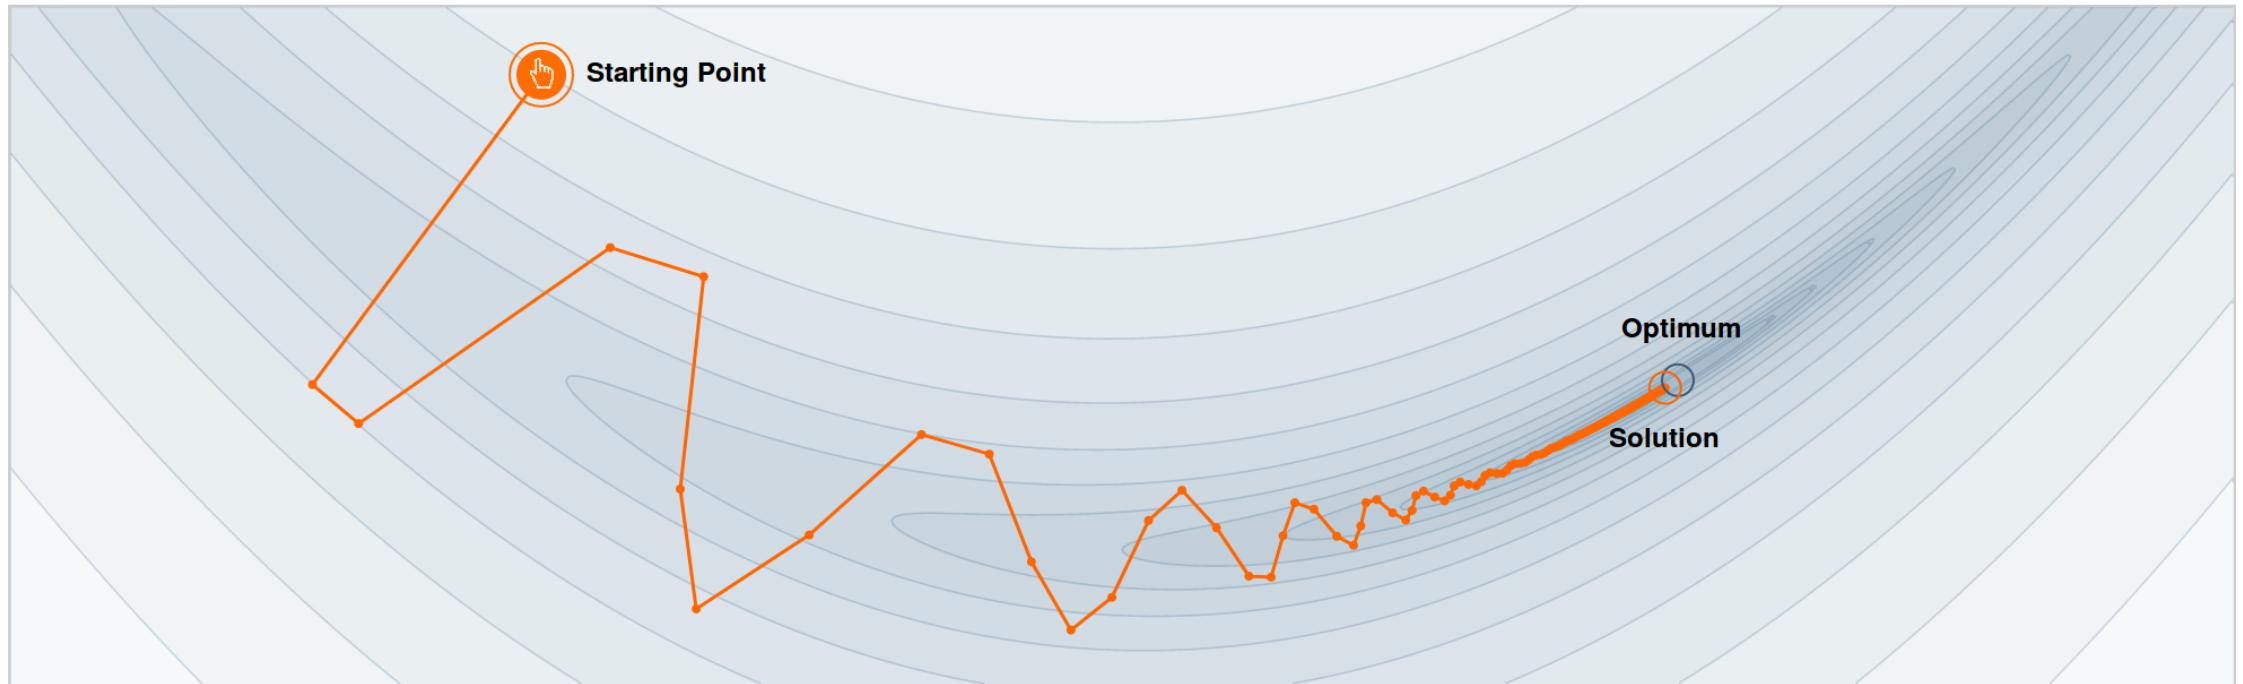


Momentum  $\beta = 0.60$

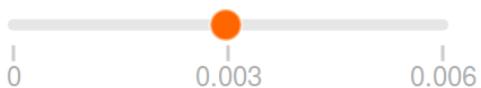


We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

## Why Momentum Really Works



Step-size  $\alpha = 0.0030$

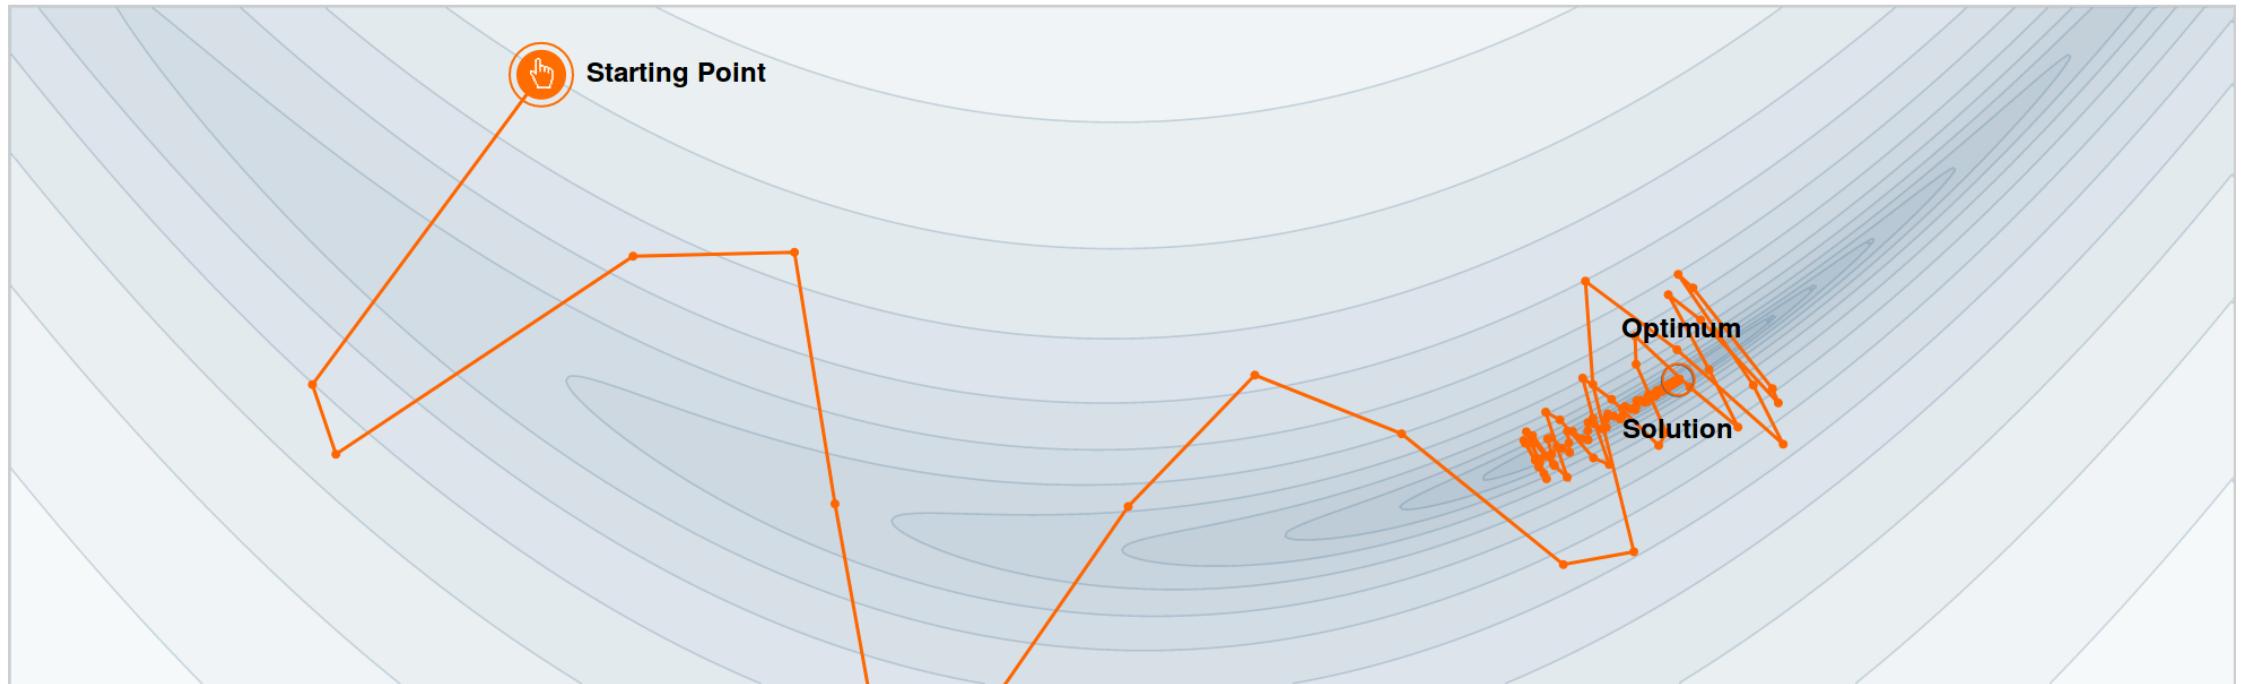


Momentum  $\beta = 0.80$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

## Why Momentum Really Works



Step-size  $\alpha = 0.0030$



Momentum  $\beta = 0.90$



We often think of Momentum as a means of dampening oscillations and speeding up the iterations, leading to faster convergence. But it has other interesting behavior. It allows a larger range of step-sizes to be used, and creates its own oscillations. What is going on?

## Why Momentum Really Works

# Alternative optimizers

- SGD (with Nesterov momentum)
  - Simple to implement
  - Very sensitive to initial value of  $\eta$
  - Need learning rate scheduling
- Adam: adaptive learning rate scale for each param
  - Global  $\eta$  set to 3e-4 often works well enough
  - Good default choice of optimizer (often)
- Many other promising methods:
  - RMSProp, Adagrad, Adadelta, Nadam, ...
  - Often takes some experimentation to find the best one



# The Karpathy Constant for Adam



Andrej Karpathy ✅

@karpathy

Following



3e-4 is the best learning rate for Adam, hands down.

4:01 AM - 24 Nov 2016

---

101 Retweets 408 Likes



23

101

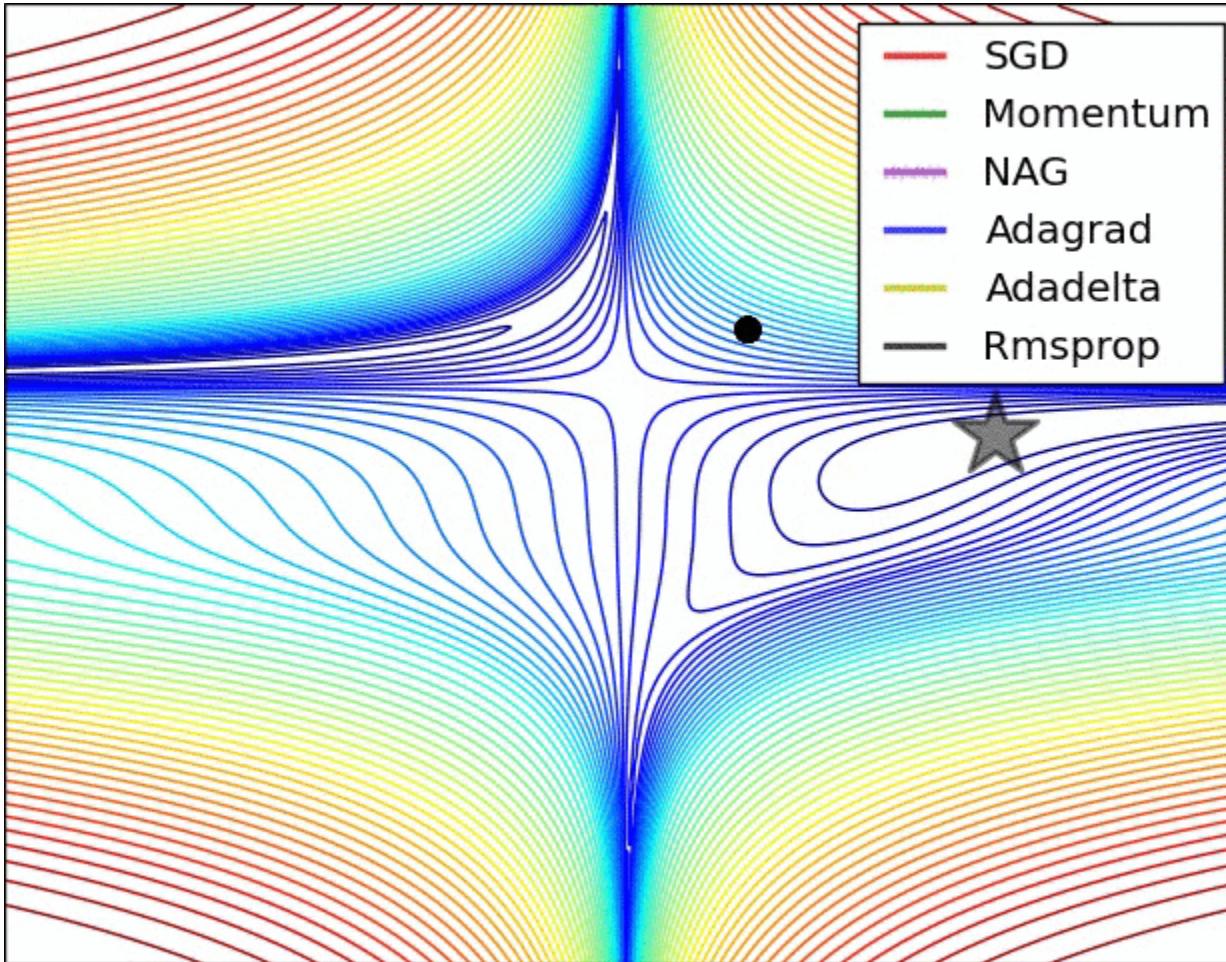


408



35

# Optimizers around a saddle point



Credits: Alec Radford

36

# Next: Lab 2!