

Deep Learning

Convolutional Neural Networks – Part II

Alex Olson

Adapted from material by Charles Ollion & Olivier Grisel


CNNs for computer Vision




[Krizhevsky 2012]



[Ciresan et al. 2013]



[Faster R-CNN - Ren 2015]



[NVIDIA dev blog]

Beyond Image Classification

CNNs

- Previous lecture: image classification

Beyond Image Classification

CNNs

- Previous lecture: image classification

Limitations

- Mostly on centered images
- Only a single object per image
- Not enough for many real world vision tasks

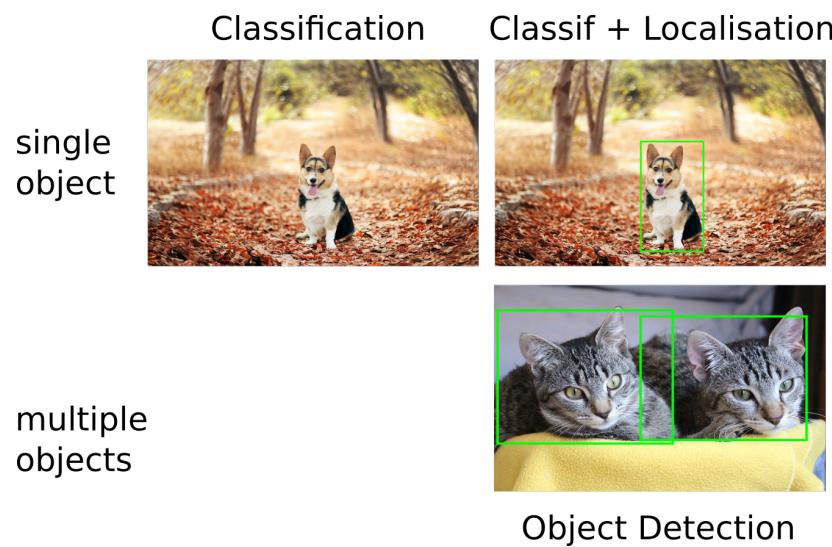
Beyond Image Classification



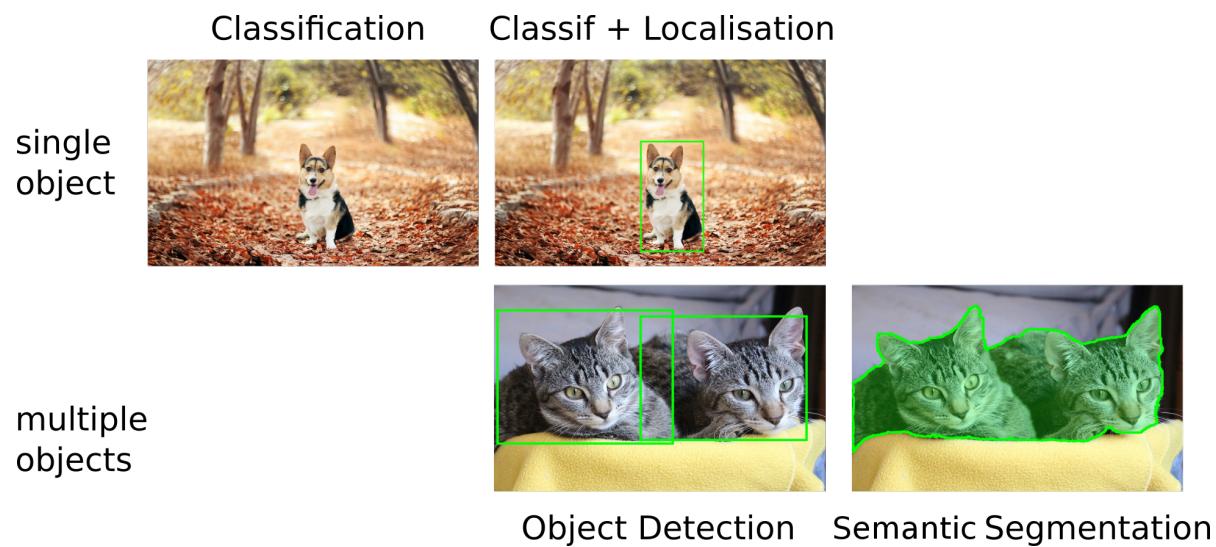
Beyond Image Classification



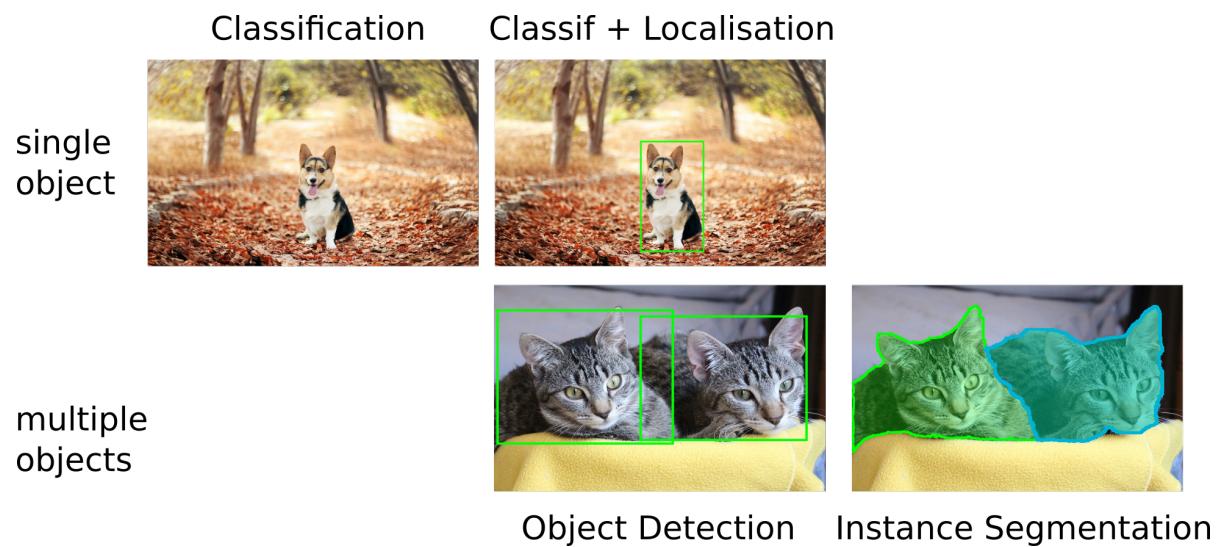
Beyond Image Classification



Beyond Image Classification



Beyond Image Classification



Outline

Simple Localisation as regression

Outline

Simple Localisation as regression

Detection Algorithms

Outline

Simple Localisation as regression

Detection Algorithms

Fully convolutional Networks

Outline

Simple Localisation as regression


Detection Algorithms

Fully convolutional Networks

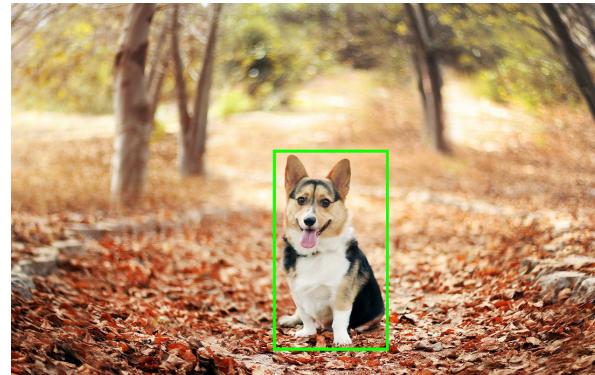
Semantic & Instance Segmentation

Localisation

Localisation

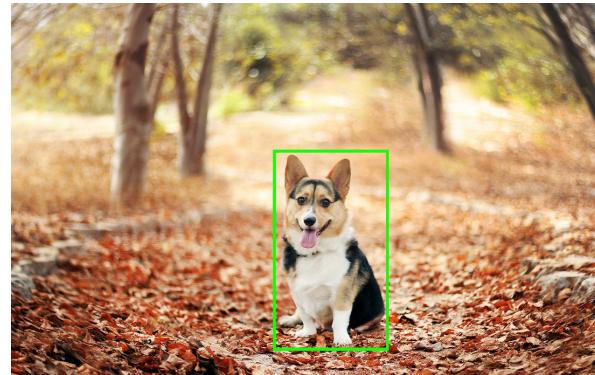


Localisation




- Single object per image
- Predict coordinates of a bounding box (x, y, w, h)

Localisation




- Single object per image
- Predict coordinates of a bounding box (x, y, w, h)
- Evaluate via Intersection over Union (IoU)

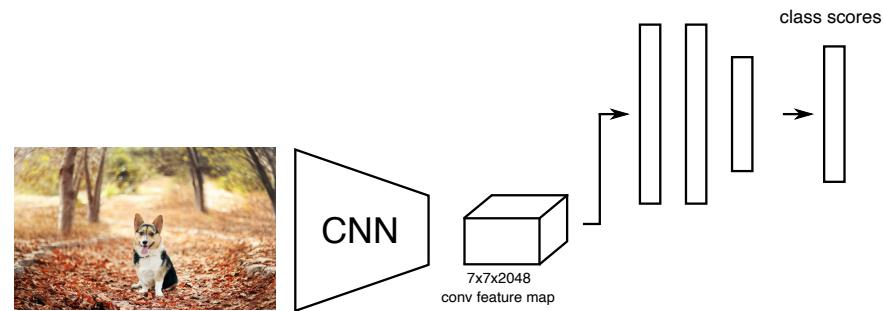
Localisation as regression



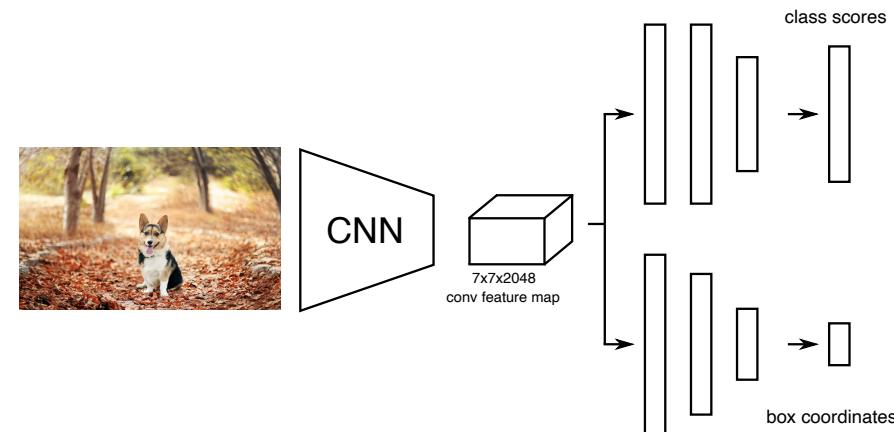
Localisation as regression



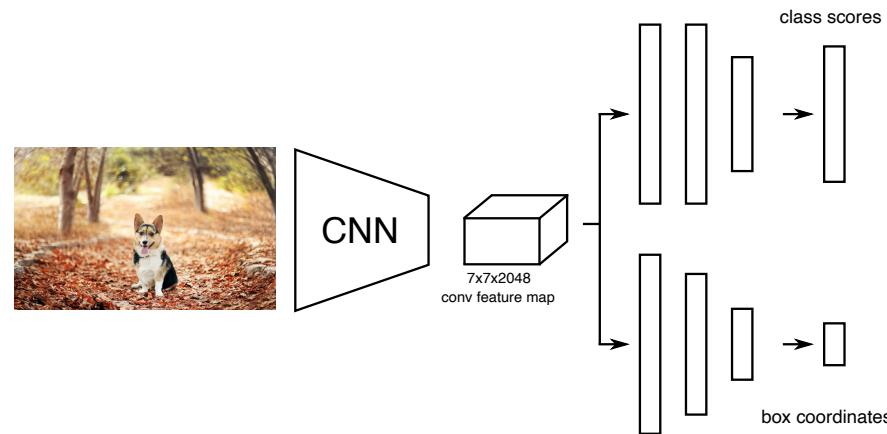
Classification + Localisation



Classification + Localisation




Classification + Localisation




- Use a pre-trained CNN on ImageNet (ex. ResNet)
- The "localisation head" is trained separately with regression

Classification + Localisation




- Use a pre-trained CNN on ImageNet (ex. ResNet)
- The "localisation head" is trained separately with regression
- Possible end-to-end finetuning of both tasks

Classification + Localisation



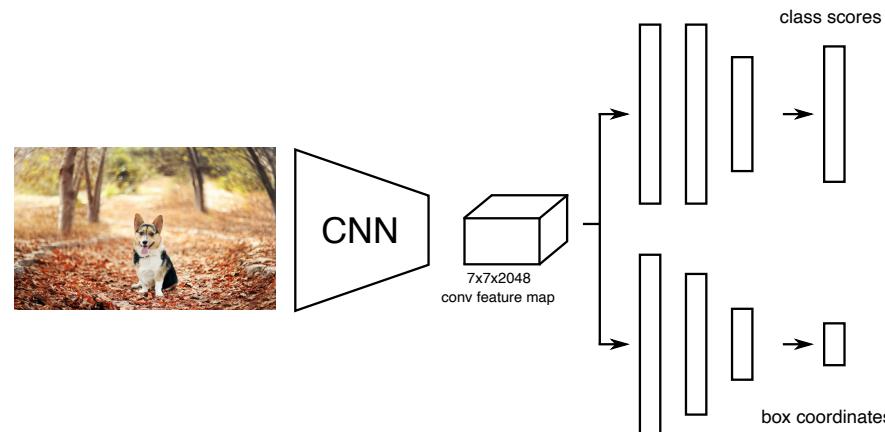
- Use a pre-trained CNN on ImageNet (ex. ResNet)
- The "localisation head" is trained separately with regression
- Possible end-to-end finetuning of both tasks
- At test time, use both heads

Classification + Localisation



C classes, 4 output dimensions (1 box)

Classification + Localisation



C classes, 4 output dimensions (1 box)

Predict exactly N objects: predict $(N \times 4)$ coordinates and $(N \times K)$ class scores

Object detection

We don't know in advance the number of objects in the image. Object detection relies on *object proposal* and *object classification*

Object proposal: find regions of interest (RoIs) in the image

Object detection

We don't know in advance the number of objects in the image. Object detection relies on *object proposal* and *object classification*

Object proposal: find regions of interest (RoIs) in the image

Object classification: classify the object in these regions

Object detection

We don't know in advance the number of objects in the image. Object detection relies on *object proposal* and *object classification*


Object proposal: find regions of interest (RoIs) in the image

Object classification: classify the object in these regions

Two main families:

- Single-Stage: A grid in the image where each cell is a proposal (SSD, YOLO, RetinaNet)
- Two-Stage: Region proposal then classification (Faster-RCNN)

YOLO



5 × 5 grid on input

Redmon, Joseph, et al. "You only look once: Unified, real-time object detection." CVPR (2016)

YOLO



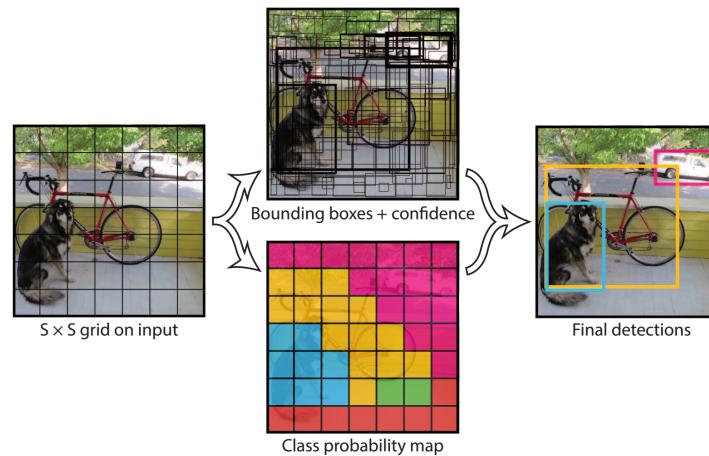
5 × 5 grid on input

For each cell of the $S \times S$ predict:

- B boxes and confidence scores C ($5 \times B$ values) + classes c

Redmon, Joseph, et al. "You only look once: Unified, real-time object detection." CVPR (2016)

YOLO




For each cell of the $S \times S$ predict:

- B boxes and confidence scores C ($5 \times B$ values) + classes c

Redmon, Joseph, et al. "You only look once: Unified, real-time object detection." CVPR (2016)

YOLO



$$\text{Final detections: } C_j * \text{prob}(c) > \text{threshold}$$

Redmon, Joseph, et al. "You only look once: Unified, real-time object detection." CVPR (2016)

YOLO

- After ImageNet pretraining, the whole network is trained end-to-end

Redmon, Joseph, et al. "You only look once: Unified, real-time object detection." CVPR (2016)

YOLO

- After ImageNet pretraining, the whole network is trained end-to-end
- The loss is a weighted sum of different regressions

$$\begin{aligned} & \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] \\ & + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[\left(\sqrt{w_i} - \sqrt{\hat{w}_i} \right)^2 + \left(\sqrt{h_i} - \sqrt{\hat{h}_i} \right)^2 \right] \\ & + \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left(C_i - \hat{C}_i \right)^2 \\ & + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} \left(C_i - \hat{C}_i \right)^2 \\ & + \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2 \quad (3) \end{aligned}$$

Redmon, Joseph, et al. "You only look once: Unified, real-time object detection." CVPR (2016)

Box Proposals

Instead of having a predefined set of box proposals, find them on the image:

- Selective Search - from pixels (not learnt, no longer used)
- Faster - RCNN - Region Proposal Network (RPN)

Box Proposals


Instead of having a predefined set of box proposals, find them on the image:

- Selective Search - from pixels (not learnt, no longer used)
- Faster - RCNN - Region Proposal Network (RPN)

Crop-and-resize operator (RoI-Pooling):


- Input: convolutional map + N regions of interest
- Output: tensor of $N \times 7 \times 7 \times$ depth boxes
- Allows to propagate gradient only on interesting regions, and efficient computation

Faster-RCNN



Ren, Shaoqing, et al. "Faster r-cnn: Towards real-time object detection with region proposal networks." NIPS 2015


Faster-RCNN



- Train jointly RPN and other head

Ren, Shaoqing, et al. "Faster r-cnn: Towards real-time object detection with region proposal networks." NIPS 2015

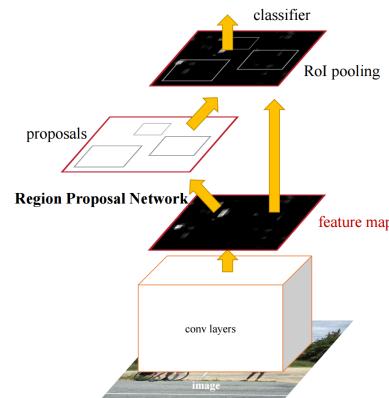
Faster-RCNN



- Train jointly RPN and other head
- 200 box proposals, gradient propagated only in positive boxes

Ren, Shaoqing, et al. "Faster r-cnn: Towards real-time object detection with region proposal networks." NIPS 2015

Faster-RCNN



- Train jointly RPN and other head
- 200 box proposals, gradient propagated only in positive boxes
- Region proposal is translation invariant, compared to YOLO

Ren, Shaoqing, et al. "Faster r-cnn: Towards real-time object detection with region proposal networks." NIPS 2015

Measuring performance

method	test size shorter edge/max size	feature pyramid	align	mAP@[0.5:0.95]	AP _s	AP _m	AP _t
R-FCN [17]	600/1000			32.1	12.8	34.9	46.1
Faster R-CNN (2fc)	600/1000			30.3	9.9	32.2	47.4
Deformable [3]	600/1000		✓	34.5	14.0	37.7	50.3
G-RMI [13]	600/1000			35.6	-	-	-
FPN [19]	800/1200	✓		36.2	18.2	39.0	48.2
Mask R-CNN [7]	800/1200	✓	✓	38.2	20.1	41.1	50.2
RetinaNet [20]	800/1200	✓		37.8	20.2	41.1	49.2
RetinaNet ms-train [20]	800/1200	✓		39.1	21.8	42.7	50.2
Light head R-CNN	800/1200		✓	39.5	21.8	43.0	50.7
Light head R-CNN ms-train	800/1200		✓	40.8	22.7	44.3	52.8
Light head R-CNN	800/1200	✓	✓	41.5	25.2	45.3	53.1

Measures: mean Average Precision mAP at given IoU thresholds

Measuring performance

method	test size shorter edge/max size	feature pyramid	align	mAP@[0.5:0.95]	AP _s	AP _m	AP _t
R-FCN [17]	600/1000			32.1	12.8	34.9	46.1
Faster R-CNN (2fc)	600/1000			30.3	9.9	32.2	47.4
Deformable [3]	600/1000		✓	34.5	14.0	37.7	50.3
G-RMI [13]	600/1000			35.6	-	-	-
FPN [19]	800/1200	✓		36.2	18.2	39.0	48.2
Mask R-CNN [7]	800/1200	✓	✓	38.2	20.1	41.1	50.2
RetinaNet [20]	800/1200	✓		37.8	20.2	41.1	49.2
RetinaNet ms-train [20]	800/1200	✓		39.1	21.8	42.7	50.2
Light head R-CNN	800/1200		✓	39.5	21.8	43.0	50.7
Light head R-CNN ms-train	800/1200		✓	40.8	22.7	44.3	52.8
Light head R-CNN	800/1200	✓	✓	41.5	25.2	45.3	53.1

Measures: mean Average Precision mAP at given IoU thresholds

- AP @0.5 for class "cat": average precision for the class, where $IoU(box^{pred}, box^{true}) > 0.5$

State-of-the-art

Model	FLOPs	# Params	AP _{val}	AP _{test-dev}
SpineNet-190 (1536) [11]	2076B	176.2M	52.2	52.5
DetectoRS ResNeXt-101-64x4d [43]	—	—	—	55.7 [†]
SpineNet-190 (1280) [11]	1885B	164M	52.6	52.8
SpineNet-190 (1280) w/ self-training [71]	1885B	164M	54.2	54.3
EfficientDet-D7x (1536) [56]	410B	77M	54.4	55.1
YOLOv4-P7 (1536) [60]	—	—	—	55.8 [†]
Cascade Eff-B7 NAS-FPN (1280)	1440B	185M	54.5	54.8
w/ Copy-Paste	1440B	185M	(+1.4) 55.9	(+1.2) 56.0
w/ self-training Copy-Paste	1440B	185M	(+2.5) 57.0	(+2.5) 57.3

State-of-the-art

Model	FLOPs	# Params	AP _{val}	AP _{test-dev}
SpineNet-190 (1536) [11]	2076B	176.2M	52.2	52.5
DetectoRS ResNeXt-101-64x4d [43]	—	—	—	55.7 [†]
SpineNet-190 (1280) [11]	1885B	164M	52.6	52.8
SpineNet-190 (1280) w/ self-training [71]	1885B	164M	54.2	54.3
EfficientDet-D7x (1536) [56]	410B	77M	54.4	55.1
YOLOv4-P7 (1536) [60]	—	—	—	55.8 [†]
Cascade Eff-B7 NAS-FPN (1280)	1440B	185M	54.5	54.8
w/ Copy-Paste	1440B	185M	(+1.4) 55.9	(+1.2) 56.0
w/ self-training Copy-Paste	1440B	185M	(+2.5) 57.0	(+2.5) 57.3

- Larger image sizes, larger and better models, better augmented data
- <https://paperswithcode.com/sota/object-detection-on-coco>

Segmentation

Segmentation

Output a class map for each pixel (here: dog vs background)



Segmentation

Output a class map for each pixel (here: dog vs background)



- **Instance segmentation:** specify each object instance as well (two dogs have different instances)


Segmentation

Output a class map for each pixel (here: dog vs background)



- **Instance segmentation:** specify each object instance as well (two dogs have different instances)
- This can be done through **object detection + segmentation**


Convolutionize



- Slide the network with an input of $(224, 224)$ over a larger image. Output of varying spatial size

Long, Jonathan, et al. "Fully convolutional networks for semantic segmentation." CVPR 2015


Convolutionize



- Slide the network with an input of $(224, 224)$ over a larger image. Output of varying spatial size
- **Convolutionize:** change Dense $(4096, 1000)$ to 1×1 Convolution, with 4096, 1000 input and output channels

Long, Jonathan, et al. "Fully convolutional networks for semantic segmentation." CVPR 2015


Convolutionize



- Slide the network with an input of $(224, 224)$ over a larger image. Output of varying spatial size
- **Convolutionize:** change Dense $(4096, 1000)$ to 1×1 Convolution, with 4096, 1000 input and output channels
- Gives a coarse segmentation (no extra supervision)

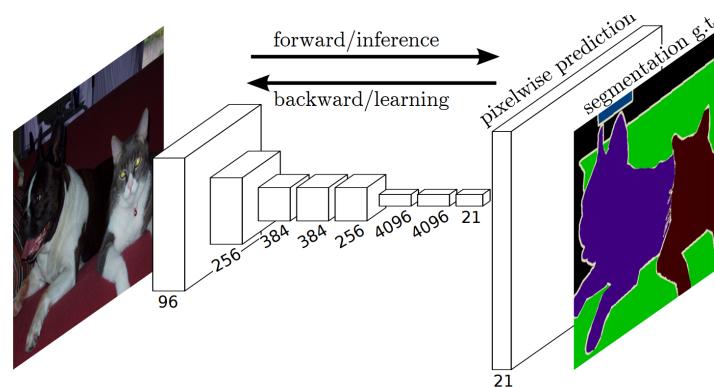
Long, Jonathan, et al. "Fully convolutional networks for semantic segmentation." CVPR 2015

Fully Convolutional Network



Long, Jonathan, et al. "Fully convolutional networks for semantic segmentation." CVPR 2015


Fully Convolutional Network



- Predict / backpropagate for every output pixel

Long, Jonathan, et al. "Fully convolutional networks for semantic segmentation." CVPR 2015


Fully Convolutional Network



- Predict / backpropagate for every output pixel
- Aggregate maps from several convolutions at different scales for more robust results


Long, Jonathan, et al. "Fully convolutional networks for semantic segmentation." CVPR 2015

Deconvolution




Noh, Hyeonwoo, et al. "Learning deconvolution network for semantic segmentation." ICCV 2015

Deconvolution




- "Deconvolution": transposed convolutions



Noh, Hyeonwoo, et al. "Learning deconvolution network for semantic segmentation." ICCV 2015


Deconvolution



- skip connections between corresponding convolution and deconvolution layers

Noh, Hyeonwoo, et al. "Learning deconvolution network for semantic segmentation." ICCV 2015


Deconvolution



- skip connections between corresponding convolution and deconvolution layers
- sharper masks by using precise spatial information (early layers)
- better object detection by using semantic information (late layers)


Noh, Hyeonwoo, et al. "Learning deconvolution network for semantic segmentation." ICCV 2015

Hourglass network



Newell, Alejandro, et al. "Stacked Hourglass Networks for Human Pose Estimation." ECCV 2016


Hourglass network



- U-Net like architectures repeated sequentially

Newell, Alejandro, et al. "Stacked Hourglass Networks for Human Pose Estimation." ECCV 2016


Hourglass network



- U-Net like architectures repeated sequentially
- Each block refines the segmentation for the following
- Each block has a segmentation loss


Newell, Alejandro, et al. "Stacked Hourglass Networks for Human Pose Estimation." ECCV 2016

Mask-RCNN



K. He and al. Mask Region-based Convolutional Network (Mask R-CNN) NIPS 2017


Mask-RCNN



Faster-RCNN architecture with a third, binary mask head


K. He and al. Mask Region-based Convolutional Network (Mask R-CNN) NIPS 2017

Results



K. He and al. Mask Region-based Convolutional Network (Mask R-CNN) NIPS 2017

Results



- Mask results are still coarse (low mask resolution)
- Excellent instance generalization

K. He and al. Mask Region-based Convolutional Network (Mask R-CNN) NIPS 2017

Results



He, Kaiming, et al. "Mask r-cnn." Internal Conference on Computer Vision (ICCV), 2017.

State-of-the-art & links

Most benchmarks and recent architectures are reported here:

<https://paperswithcode.com/area/computer-vision>

State-of-the-art & links

Most benchmarks and recent architectures are reported here:

<https://paperswithcode.com/area/computer-vision>

Tensorflow

[object detection API](#)

State-of-the-art & links

Most benchmarks and recent architectures are reported here:

<https://paperswithcode.com/area/computer-vision>

Tensorflow

[object detection API](#)

Pytorch

Detectron <https://github.com/facebookresearch/Detectron>

- Mask-RCNN, Retina Net and other architectures
- Focal loss, Feature Pyramid Networks, etc.

Take away NN for Vision

Pre-trained features as a basis

- ImageNet: centered objects, very broad image domain
- 1M+ labels and many different classes resulting in very general and disentangling representations
- Better Networks (i.e. ResNet vs VGG) have a huge impact

Take away NN for Vision

Pre-trained features as a basis

- ImageNet: centered objects, very broad image domain
- 1M+ labels and many different classes resulting in very general and disentangling representations
- Better Networks (i.e. ResNet vs VGG) have a huge impact

Fine tuning

- Add new layers on top of convolutional or dense layer of CNNs
- Fine tune the whole architecture end-to-end
- Make use of a smaller dataset but with richer labels (bounding boxes, masks...)

Next: Lab 5!